

Multivariate Linear Regression on ENB Dataset

```
In [1]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LinearRegression
        from sklearn.metrics import mean_squared_error, r2_score
        import matplotlib.pyplot as plt
```


Load dataset

```
In [2]: df = pd.read_csv("../dataset/energy_efficiency_data.csv")
```

```
In [3]: # print first 5 rows
        df.head()
```

```
Out[3]:
```

| | Relative_Compactness | Surface_Area | Wall_Area | Roof_Area | Overall_Height | Orientation |
|---|----------------------|--------------|-----------|-----------|----------------|-------------|
| 0 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | |
| 1 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | |
| 2 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | |
| 3 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | |
| 4 | 0.90 | 563.5 | 318.5 | 122.50 | 7.0 | |



```
In [4]: # check for missing values
        df.isnull().sum()
```

```
Out[4]: Relative_Compactness      0
        Surface_Area              0
        Wall_Area                 0
        Roof_Area                 0
        Overall_Height            0
        Orientation               0
        Glazing_Area              0
        Glazing_Area_Distribution  0
        Heating_Load              0
        Cooling_Load              0
        dtype: int64
```

```
In [5]: df.describe()
```

Out[5]:

| | Relative_Compactness | Surface_Area | Wall_Area | Roof_Area | Overall_Height | Ori |
|--------------|----------------------|--------------|------------|------------|----------------|-----|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768 |
| mean | 0.764167 | 671.708333 | 318.500000 | 176.604167 | 5.250000 | 3 |
| std | 0.105777 | 88.086116 | 43.626481 | 45.165950 | 1.75114 | 1 |
| min | 0.620000 | 514.500000 | 245.000000 | 110.250000 | 3.500000 | 2 |
| 25% | 0.682500 | 606.375000 | 294.000000 | 140.875000 | 3.500000 | 2 |
| 50% | 0.750000 | 673.750000 | 318.500000 | 183.750000 | 5.250000 | 3 |
| 75% | 0.830000 | 741.125000 | 343.000000 | 220.500000 | 7.000000 | 4 |
| max | 0.980000 | 808.500000 | 416.500000 | 220.500000 | 7.000000 | 5 |



In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Relative_Compactness                  768 non-null    float64
1   Surface_Area                         768 non-null    float64
2   Wall_Area                           768 non-null    float64
3   Roof_Area                           768 non-null    float64
4   Overall_Height                       768 non-null    float64
5   Orientation                          768 non-null    int64
6   Glazing_Area                        768 non-null    float64
7   Glazing_Area_Distribution            768 non-null    int64
8   Heating_Load                        768 non-null    float64
9   Cooling_Load                        768 non-null    float64
dtypes: float64(8), int64(2)
memory usage: 60.1 KB
```

Select Features (X) and Target (y)

In [7]: `# Dataset has 8 features (columns 0 to 7) and 2 target variables (last 2 columns)`
`X = df.iloc[:, :-2] # all rows, first 8 columns → features`
`y = df.iloc[:, 8:] # all rows, last 2 columns → targets (Heating, Cooling Load)`

In [10]: `print("The shape of X (Features): ", X.shape)`
`print("The shape of y (target variables): ", y.shape)`

The shape of X (Features): (768, 8)
The shape of y (target variables): (768, 2)

Split dataset into train and test samples

In [11]: `# 80% training and 20% testing`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)`

In [12]: `X_train.shape`

Out[12]: (614, 8)

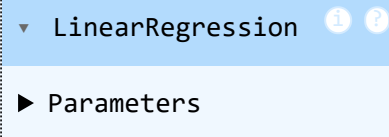
In [13]: `X_test.shape`

Out[13]: (154, 8)

Train Multi-variate Linear Regression Model

In [14]: `model = LinearRegression()`

In [15]: `model.fit(X_train, y_train)`

Out[15]: 

Make Predictions on test sample

In [16]: `y_pred = model.predict(X_test)`

Evaluate the model performance

In [17]:

```
# Calculate error metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("\nMean Squared Error (MSE):", mse)
print("R2 Score:", r2)
```

Mean Squared Error (MSE): 9.523307751573416
R² Score: 0.9027048110077098

Compare Predictions with Actual Values

In [19]:

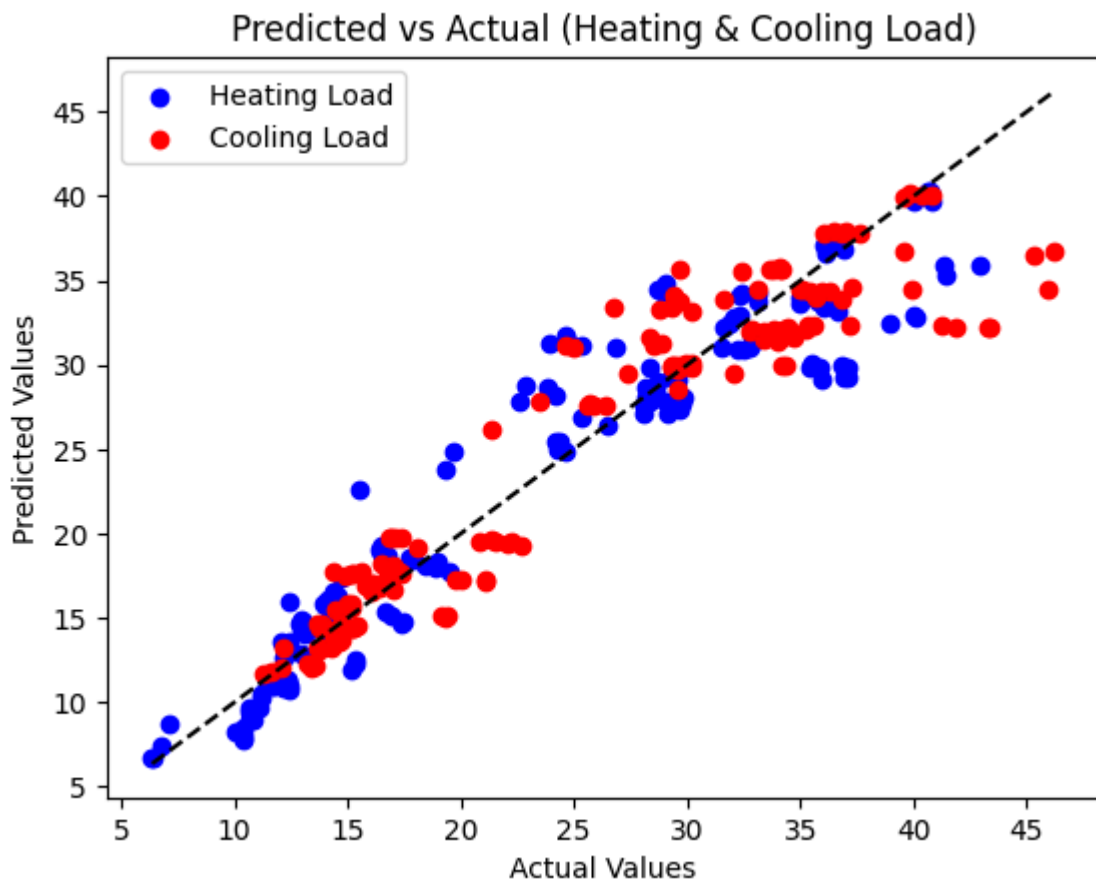
```
print("\nFirst 5 Predictions:\n", y_pred[:5])
print("\nFirst 5 Actual Values:\n", y_test.values[:5])
```

First 5 Predictions:
[[18.86296033 19.75455641]
[14.04938012 16.77160938]
[31.31560297 32.02497722]
[35.90050966 36.70240832]
[15.33519734 17.27670437]]

First 5 Actual Values:
[[16.47 16.9]
[13.17 16.39]
[32.82 32.78]
[41.32 46.23]
[16.69 19.76]]

Visualization

```
In [20]: # Scatter plot: Predicted vs Actual for Heating Load (first target)
plt.scatter(y_test.iloc[:,0], y_pred[:,0], color="blue", label="Heating Load")
plt.scatter(y_test.iloc[:,1], y_pred[:,1], color="red", label="Cooling Load")
plt.plot([y_test.min().min(), y_test.max().max()],
         [y_test.min().min(), y_test.max().max()],
         color="black", linestyle="--") # perfect prediction line
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Predicted vs Actual (Heating & Cooling Load)")
plt.legend()
plt.show()
```



```
In [24]: import numpy as np
# Loop through each target separately
target_names = ["Heating Load", "Cooling Load"]

for i, name in enumerate(target_names):
    mse_i = mean_squared_error(y_test.iloc[:, i], y_pred[:, i])
    r2_i = r2_score(y_test.iloc[:, i], y_pred[:, i])
    print(f"\n{name}:")
    print(f"    MSE: {mse_i:.2f}")
    print(f"    RMSE: {np.sqrt(mse_i):.2f}")
    print(f"    R²: {r2_i:.4f}")
```

Heating Load:

MSE: 9.15

RMSE: 3.03

R^2 : 0.9122

Cooling Load:

MSE: 9.89

RMSE: 3.15

R^2 : 0.8932

```
In [22]: plt.figure(figsize=(8,6))

# Scatter plots for actual vs predicted
plt.scatter(y_test.iloc[:,0], y_pred[:,0], color="blue", alpha=0.6, label="Heating Load")
plt.scatter(y_test.iloc[:,1], y_pred[:,1], color="red", alpha=0.6, label="Cooling Load")

# Heating Load perfect line
plt.plot([y_test.iloc[:,0].min(), y_test.iloc[:,0].max()],
         [y_test.iloc[:,0].min(), y_test.iloc[:,0].max()],
         color="blue", linestyle="--", label="Perfect Heating")

# Cooling Load perfect line
plt.plot([y_test.iloc[:,1].min(), y_test.iloc[:,1].max()],
         [y_test.iloc[:,1].min(), y_test.iloc[:,1].max()],
         color="red", linestyle="--", label="Perfect Cooling")

# Labels and title
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Predicted vs Actual (Heating & Cooling Load)")
plt.legend()
plt.show()
```

Predicted vs Actual (Heating & Cooling Load)

